

Deep Learning

Spring 2022

Prof. Gilles Louppe
g.louppe@uliege.be

Philosophy

Thorough and detailed

- Understand the foundations and the landscape of deep learning.
- Be able to write from scratch, debug and run (some) deep learning algorithms.

State-of-the-art

- Introduction to materials new from research (\leq 5 years old).
- Understand some of the open questions and challenges in the field.

Practical

- Fun and challenging course project.

This course is given by:


- Theory: Gilles Louppe (g.louppe@uliege.be)
- Projects and guidance:
 - Arnaud Delaunoy
 - François Rozet
 - Gaspard Lambrechts
 - Yann Claes
 - Antoine Wehenkel



Course hub

All important information about the course is available on the course web page github.com/glouppe/info8010-deep-learning:

- Schedule
- Slides and materials
- Homeworks and project

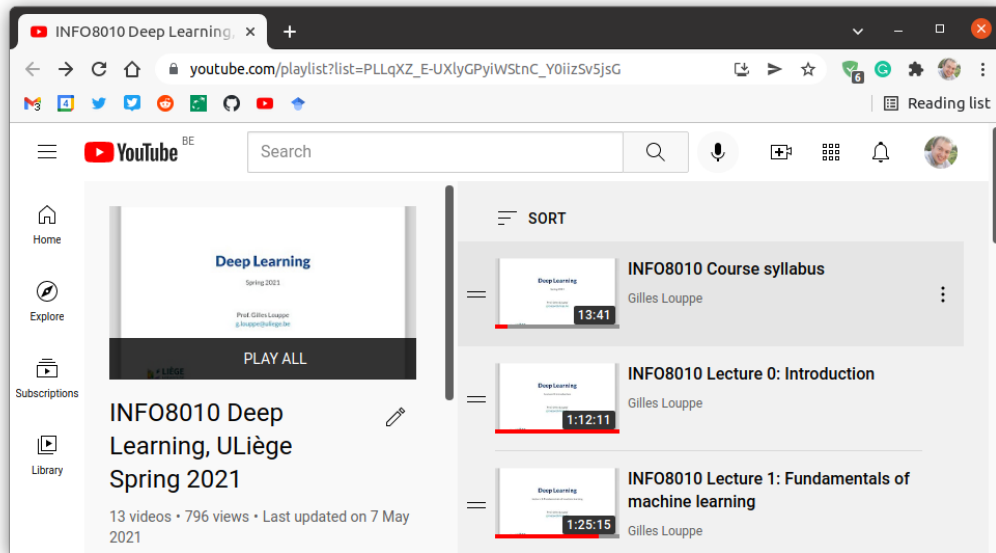


The screenshot shows a web browser displaying the GitHub repository page for 'glouppe/info8010-deep-learning'. The page title is 'INFO8010 - Deep Learning'. Below the title, it states 'Lectures for INFO8010 - Deep Learning, ULiège, Spring 2022.' and lists the instructor (Gilles Louppe), teaching assistants (Arnaud Delaunoy, François Rozet, Antoine Wehenkel), the time (Spring 2022, Friday 8:30 AM), and the classroom (R3 / B28). The 'Agenda' section contains a table with dates and topics.

Date	Topic
February 11	Course syllabus [PDF] [video] Lecture 0: Introduction [PDF] [video] Lecture 1: Fundamentals of machine learning [PDF] [Side notes] [video]
February 18	Lecture 2: Multi-layer perceptron [PDF] [Side notes] [video]
February 25	Lecture 3: Automatic differentiation [PDF] [Side notes] [video]

Videos


Videos from Spring 2021 are available at <https://bit.ly/3roQmea>.



Textbook

None!

But we would recommend "Dive into Deep Learning" (d2l.ai) for a comprehensive and practical introduction to the field.



Projects

Reading assignment

Read, summarize and criticize a major scientific paper in deep learning. Details to be announced soon.

ARTICLE

doi:10.1038/nature16961

Mastering the game of Go with deep neural networks and tree search

David Silver^{1*}, Aja Huang^{1*}, Chris J. Maddison¹, Arthur Guez¹, Laurent Sifre¹, George van den Driessche¹, Julian Schrittwieser¹, Ioannis Antonoglou¹, Veda Panneershelvam¹, Marc Lanctot¹, Sander Dieleman¹, Dominik Grewe¹, John Nham², Nal Kalchbrenner¹, Ilya Sutskever², Timothy Lillicrap¹, Madeleine Leach¹, Koray Kavukcuoglu¹, Thore Graepel¹ & Demis Hassabis¹

The game of Go has long been viewed as the most challenging of classic games for artificial intelligence owing to its enormous search space and the difficulty of evaluating board positions and moves. Here we introduce a new approach to computer Go that uses 'value networks' to evaluate board positions and 'policy networks' to select moves. These deep neural networks are trained by a novel combination of supervised learning from human expert games, and reinforcement learning from games of self-play. Without any lookahead search, the neural networks play Go at the level of state-of-the-art Monte Carlo tree search programs that simulate thousands of random games of self-play. We also introduce a new search algorithm that combines Monte Carlo simulation with value and policy networks. Using this search algorithm, our program AlphaGo achieved a 99.8% winning rate against other Go programs, and defeated the human European Go champion by 5 games to 0. This is the first time that a computer program has defeated a human professional player in the full-sized game of Go, a feat previously thought to be at least a decade away.

All games of perfect information have an optimal value function, $v^*(s)$, which determines the outcome of the game, from every board position or state s , under perfect play by all players. These games may be solved by recursively computing the optimal value function in a search tree containing approximately b^d possible sequences of moves, where b is the game's breadth (number of legal moves per position) and d is its depth (game length). In large games, such as chess ($b \approx 35$, $d \approx 80$)¹ and especially Go ($b \approx 250$, $d \approx 150$)², exhaustive search is infeasible^{3,4}, but the effective search space can be reduced by two general principles. First, the depth of the search may be reduced by position evaluation: truncating the search tree at state s and replacing the subtree below s by an approximate value function $v(s) \approx v^*(s)$ that predicts the outcome from state s . This approach has led to superhuman performance in chess⁵, checkers⁶ and othello⁶, but it was believed to be intractable in Go due to the complexity of the game⁷. Second, the breadth of the search may be reduced by sampling actions from a policy $p(a|s)$ that is a probability distribution over possible moves a in position s . For example, Monte Carlo rollouts⁸ search to maximum depth without branching at all, by sampling long sequences of actions for both players from a policy p . Averaging over such rollouts can provide an effective position evaluation, achieving superhuman performance in backgammon⁹ and

policies^{13–15} or value functions¹⁶ based on a linear combination of input features.

Recently, deep convolutional neural networks have achieved unprecedented performance in visual domains: for example, image classification¹⁷, face recognition¹⁸, and playing Atari games¹⁹. They use many layers of neurons, each arranged in overlapping tiles, to construct increasingly abstract, localized representations of an image²⁰. We employ a similar architecture for the game of Go. We pass in the board position as a 19×19 image and use convolutional layers to construct a representation of the position. We use these neural networks to reduce the effective depth and breadth of the search tree: evaluating positions using a value network, and sampling actions using a policy network.

We train the neural networks using a pipeline consisting of several stages of machine learning (Fig. 1). We begin by training a supervised learning (SL) policy network p_π directly from expert human moves. This provides fast, efficient learning updates with immediate feedback and high-quality gradients. Similar to prior work^{13,15}, we also train a fast policy p_π that can rapidly sample actions during rollouts. Next, we train a reinforcement learning (RL) policy network p_π that improves the SL policy network by optimizing the final outcome of games of self-play. This adjusts the policy towards the correct goal of winning games,

Homeworks

Short exercises to get you started with the practicals of deep learning.

Project

Project of your choosing. Details to be announced soon.

Evaluation

- Oral exam (50%)
- Reading assignment (10%)
- Projects (40%)
 - Homeworks (10%) (optional)
 - Programming project (30% or 40%)

The reading assignment and the projects are **mandatory** for presenting the exam.

